

```
In [1]: 1 import pandas as pd  
        2
```

```
In [2]: 1 data=pd.read_csv("/home/placement/Downloads/Titanic Dataset.csv")
```

```
In [3]: 1 data.describe()
```

Out[3]:

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

In [4]: 1 data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null    int64
1   Survived        891 non-null    int64
2   Pclass          891 non-null    int64
3   Name            891 non-null    object
4   Sex             891 non-null    object
5   Age            714 non-null    float64
6   SibSp          891 non-null    int64
7   Parch          891 non-null    int64
8   Ticket          891 non-null    object
9   Fare           891 non-null    float64
10  Cabin          204 non-null    object
11  Embarked       889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

In [5]: 1 list(data)

```
Out[5]: ['PassengerId',
'Survived',
'Pclass',
'Name',
'Sex',
'Age',
'SibSp',
'Parch',
'Ticket',
'Fare',
'Cabin',
'Embarked']
```

```
In [6]: 1 data.isna().sum()#no.of null values
```

```
Out[6]: PassengerId      0  
Survived      0  
Pclass        0  
Name          0  
Sex           0  
Age          177  
SibSp         0  
Parch         0  
Ticket        0  
Fare          0  
Cabin        687  
Embarked      2  
dtype: int64
```

```
In [7]: 1 data1=data.drop(['Cabin', 'Name', 'PassengerId', 'Ticket', 'SibSp', 'Parch'],axis=1)#deleting columns
```

In [8]: 1 data1

Out[8]:

	Survived	Pclass	Sex	Age	Fare	Embarked
0	0	3	male	22.0	7.2500	S
1	1	1	female	38.0	71.2833	C
2	1	3	female	26.0	7.9250	S
3	1	1	female	35.0	53.1000	S
4	0	3	male	35.0	8.0500	S
...
886	0	2	male	27.0	13.0000	S
887	1	1	female	19.0	30.0000	S
888	0	3	female	NaN	23.4500	S
889	1	1	male	26.0	30.0000	C
890	0	3	male	32.0	7.7500	Q

891 rows × 6 columns

In [9]: 1 data1.isna().sum()*#finding null values*

Out[9]: Survived 0
Pclass 0
Sex 0
Age 177
Fare 0
Embarked 2
dtype: int64

In [10]: 1 data1.shape*#no of rows and columns*

Out[10]: (891, 6)

```
In [11]: 1 data1['Sex']=data1['Sex'].map({'male':1,'female':0})#assinging 1 to male and 0 to female using map
```

```
In [12]: 1 data1
```

```
Out[12]:
```

	Survived	Pclass	Sex	Age	Fare	Embarked
0	0	3	1	22.0	7.2500	S
1	1	1	0	38.0	71.2833	C
2	1	3	0	26.0	7.9250	S
3	1	1	0	35.0	53.1000	S
4	0	3	1	35.0	8.0500	S
...
886	0	2	1	27.0	13.0000	S
887	1	1	0	19.0	30.0000	S
888	0	3	0	NaN	23.4500	S
889	1	1	1	26.0	30.0000	C
890	0	3	1	32.0	7.7500	Q

891 rows × 6 columns

```
In [13]: 1 data1=data1.fillna(data1.mean())#fill null values using mean
          2
```

/snap/jupyter/6/lib/python3.7/site-packages/ipykernel_launcher.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

"""Entry point for launching an IPython kernel.

```
In [14]: 1 data1
```

```
Out[14]:
```

	Survived	Pclass	Sex	Age	Fare	Embarked
0	0	3	1	22.000000	7.2500	S
1	1	1	0	38.000000	71.2833	C
2	1	3	0	26.000000	7.9250	S
3	1	1	0	35.000000	53.1000	S
4	0	3	1	35.000000	8.0500	S
...
886	0	2	1	27.000000	13.0000	S
887	1	1	0	19.000000	30.0000	S
888	0	3	0	29.699118	23.4500	S
889	1	1	1	26.000000	30.0000	C
890	0	3	1	32.000000	7.7500	Q

891 rows × 6 columns

```
In [15]: 1 data1.isna().sum()
```

```
Out[15]: Survived    0
Pclass      0
Sex         0
Age         0
Fare        0
Embarked    2
dtype: int64
```

```
In [16]: 1 data2=pd.get_dummies(data1)#where the pclass it shows"1" other pclass it shows "0"
          2 data2
```

Out[16]:

	Survived	Pclass	Sex	Age	Fare	Embarked_C	Embarked_Q	Embarked_S
0	0	3	1	22.000000	7.2500	0	0	1
1	1	1	0	38.000000	71.2833	1	0	0
2	1	3	0	26.000000	7.9250	0	0	1
3	1	1	0	35.000000	53.1000	0	0	1
4	0	3	1	35.000000	8.0500	0	0	1
...
886	0	2	1	27.000000	13.0000	0	0	1
887	1	1	0	19.000000	30.0000	0	0	1
888	0	3	0	29.699118	23.4500	0	0	1
889	1	1	1	26.000000	30.0000	1	0	0
890	0	3	1	32.000000	7.7500	0	1	0

891 rows × 8 columns

```
In [17]: 1 x=data2.drop(['Survived'],axis=1)#deleting churn
```

In [18]:

```
1 x
```

Out[18]:

	Pclass	Sex	Age	Fare	Embarked_C	Embarked_Q	Embarked_S
0	3	1	22.000000	7.2500	0	0	1
1	1	0	38.000000	71.2833	1	0	0
2	3	0	26.000000	7.9250	0	0	1
3	1	0	35.000000	53.1000	0	0	1
4	3	1	35.000000	8.0500	0	0	1
...
886	2	1	27.000000	13.0000	0	0	1
887	1	0	19.000000	30.0000	0	0	1
888	3	0	29.699118	23.4500	0	0	1
889	1	1	26.000000	30.0000	1	0	0
890	3	1	32.000000	7.7500	0	1	0

891 rows × 7 columns

In [20]:

```
1 y=data['Survived']
```

In [21]:

```
1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```


In [22]: 1 x.head(5)

Out[22]:

	Pclass	Sex	Age	Fare	Embarked_C	Embarked_Q	Embarked_S
0	3	1	22.0	7.2500	0	0	1
1	1	0	38.0	71.2833	1	0	0
2	3	0	26.0	7.9250	0	0	1
3	1	0	35.0	53.1000	0	0	1
4	3	1	35.0	8.0500	0	0	1

In [23]:

```

1 from sklearn.model_selection import GridSearchCV #GridSearchCV is for parameter tuning
2 from sklearn.ensemble import RandomForestClassifier
3 cls=RandomForestClassifier()
4 n_estimators=[25,50,75,100,125,150,175,200] #number of decision trees in the forest, default = 100
5 criterion=['gini','entropy'] #criteria for choosing nodes default = 'gini'
6 max_depth=[3,5,10] #maximum number of nodes in a tree default = None (it will go till all possible nodes)
7 parameters={'n_estimators': n_estimators, 'criterion':criterion, 'max_depth':max_depth} #this will undergo grid search
8 RFC_cls = GridSearchCV(cls, parameters)
9 RFC_cls.fit(x_train,y_train)

```

Out[23]: GridSearchCV(estimator=RandomForestClassifier(),
 param_grid={'criterion': ['gini', 'entropy'],
 'max_depth': [3, 5, 10],
 'n_estimators': [25, 50, 75, 100, 125, 150, 175, 200]})

In [24]: 1 RFC_cls.best_params_

Out[24]: {'criterion': 'gini', 'max_depth': 5, 'n_estimators': 200}

In [25]: 1 cls=RandomForestClassifier(n_estimators=175,criterion='entropy',max_depth=10)

```
In [26]: 1 cls.fit(x_train,y_train)
```

```
Out[26]: RandomForestClassifier(criterion='entropy', max_depth=10, n_estimators=175)
```

```
In [27]: 1 rfy_pred=cls.predict(x_test)
```

```
In [28]: 1 rfy_pred
```

```
Out[28]: array([0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1,
                0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1,
                0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0,
                1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0,
                0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1,
                0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0,
                0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0,
                1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0,
                0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1,
                0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0,
                0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0,
                1, 0, 1, 1, 0, 0, 1, 1, 0])
```

```
In [29]: 1 from sklearn.metrics import confusion_matrix
          2 confusion_matrix(y_test,rfy_pred)
```

```
Out[29]: array([[149, 26],
                [ 35, 85]])
```

```
In [30]: 1 from sklearn.metrics import accuracy_score
          2 accuracy_score(y_test,rfy_pred)#EFFICENCY OF THE CONFUSION MATRIX
```

```
Out[30]: 0.7932203389830509
```

```
In [ ]: 1
```

